

MNIST_for_beginners_no_convolution

December 21, 2016

1 MNIST For ML Beginners

A very simple MNIST classifier. See extensive documentation at <http://tensorflow.org/tutorials/mnist/beginners/index.md>

```
In [1]: from __future__ import absolute_import
        from __future__ import division
        from __future__ import print_function

        import os.path

        import argparse
        import sys

        from tensorflow.examples.tutorials.mnist import input_data

        import tensorflow as tf

        flags = tf.app.flags
        FLAGS = flags.FLAGS
        flags.DEFINE_string('data_dir', './', 'Directory to put the training data.')

In [2]: def main(_):
        # Import data
        mnist = input_data.read_data_sets(FLAGS.data_dir, one_hot=True)

        # Create the model
        x = tf.placeholder(tf.float32, [None, 784])
        W = tf.Variable(tf.zeros([784, 10]))
        b = tf.Variable(tf.zeros([10]))
        y = tf.matmul(x, W) + b

        # Define loss and optimizer
        y_ = tf.placeholder(tf.float32, [None, 10])

        # The raw formulation of cross-entropy,
        #
```

```

# tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(tf.nn.softmax(y)),
#                               reduction_indices=[1]))
#
# can be numerically unstable.
#
# So here we use tf.nn.softmax_cross_entropy_with_logits on the raw
# outputs of 'y', and then average across the batch.
cross_entropy = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(y,
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)

#sess = tf.InteractiveSession()
with tf.Session() as sess:
    #tf.global_variables_initializer().run()
    init = tf.initialize_all_variables()
    sess.run(init)
    # Train
    for _ in range(1000):
        batch_xs, batch_ys = mnist.train.next_batch(100)
        sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})

    # Test trained model
    correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(y_, 1))
    accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
    print(">>> Test Accuracy:."+str(sess.run(accuracy, feed_dict={x: mnist

```

```
In [3]: main(_)
```

```

Extracting ./train-images-idx3-ubyte.gz
Extracting ./train-labels-idx1-ubyte.gz
Extracting ./t10k-images-idx3-ubyte.gz
Extracting ./t10k-labels-idx1-ubyte.gz
>>> Test Accuracy::0.9197

```

1.1 TensorBoard: Visualizing Learning

```

In [4]: def variable_summaries(var):
    """Attach a lot of summaries to a Tensor (for TensorBoard visualization).
    with tf.name_scope('summaries'):
        mean = tf.reduce_mean(var)
        #tf.summary.scalar('mean', mean)
        tf.scalar_summary(var.name+'_mean', mean)
        with tf.name_scope('stddev'):
            stddev = tf.sqrt(tf.reduce_mean(tf.square(var - mean)))
            #tf.summary.scalar('stddev', stddev)
            tf.scalar_summary(var.name+'_stddev', stddev)
            #tf.summary.scalar('max', tf.reduce_max(var))
            tf.scalar_summary(var.name+'_max', tf.reduce_max(var))

```

```

    #tf.summary.scalar('min', tf.reduce_min(var))
    tf.scalar_summary(var.name+'_min', tf.reduce_min(var))

def main2(_):
    # Import data
    mnist = input_data.read_data_sets(FLAGS.data_dir, one_hot=True)

    # Create the model
    x = tf.placeholder(tf.float32, [None, 784])

    with tf.name_scope('W'):
        W = tf.Variable(tf.zeros([784, 10]))
        variable_summaries(W)

    with tf.name_scope('b'):
        b = tf.Variable(tf.zeros([10]))
        variable_summaries(b)

    with tf.name_scope('y'):
        y = tf.matmul(x, W) + b
        variable_summaries(y)

    # Define loss and optimizer
    y_ = tf.placeholder(tf.float32, [None, 10])

    with tf.name_scope('cross_entropy'):
        cross_entropy = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits
#tf.summary.scalar('cross_entropy', cross_entropy)
        tf.scalar_summary('cross_entropy', cross_entropy)

    with tf.name_scope('train_step'):
        train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)

    # Test trained model
    with tf.name_scope('accuracy'):
        with tf.name_scope('correct_prediction'):
            correct_prediction = tf.equal(tf.argmax(y, 1), tf.argmax(y_, 1))

            with tf.name_scope('accuracy'):
                accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
#tf.summary.scalar('accuracy', accuracy)
        tf.scalar_summary('accuracy', accuracy)

    #####
    init = tf.initialize_all_variables()

    # Create a saver for writing training checkpoints.

```

```

saver = tf.train.Saver()

#sess = tf.InteractiveSession()
with tf.Session() as sess:
    #tf.global_variables_initializer().run()

    # Merge all the summaries and write them out to ./logs (by default)
    merged = tf.merge_all_summaries()
    writer = tf.train.SummaryWriter(FLAGS.data_dir + '/_logs', sess.graph)

    sess.run(init)

    # Train
    iterations = 1000
    for i in range(iterations):
        batch_xs, batch_ys = mnist.train.next_batch(100)
        sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})

        if i % 100 == 0 or i == (iterations-1):
            summary = sess.run(merged, feed_dict={x: batch_xs, y_: batch_ys})
            writer.add_summary(summary, i)
            summary, acc = sess.run([merged, accuracy], feed_dict={x: mnist.t
            writer.add_summary(summary, i)
            writer.flush()

            checkpoint_file = os.path.join(FLAGS.data_dir + '/_logs', 'checkp
            saver.save(sess, checkpoint_file, global_step=i)

            print('>>> Test Accuracy [%s/%s]: %s' % (i, iterations, acc))

```

In [5]: main2(_)

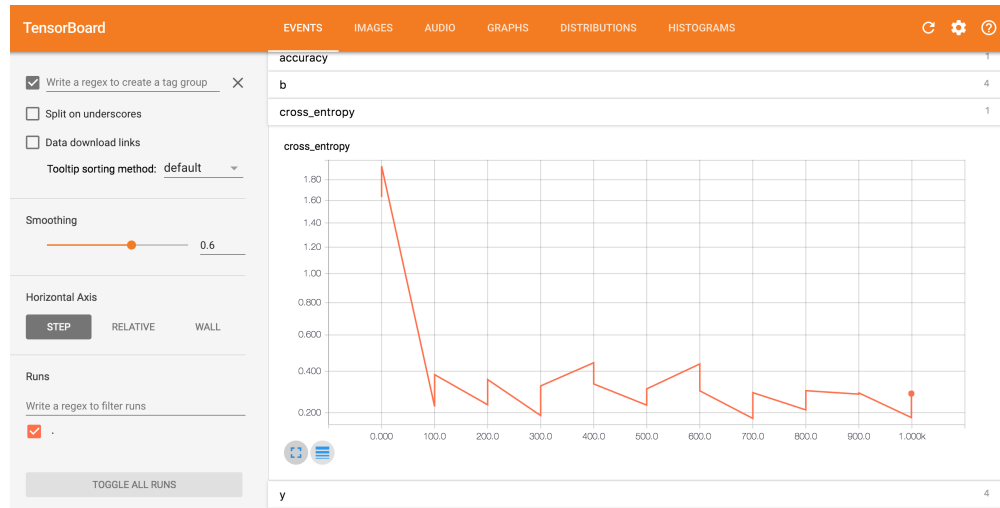
```

Extracting ./train-images-idx3-ubyte.gz
Extracting ./train-labels-idx1-ubyte.gz
Extracting ./t10k-images-idx3-ubyte.gz
Extracting ./t10k-labels-idx1-ubyte.gz
>>> Test Accuracy [0/1000]: 0.4075
>>> Test Accuracy [100/1000]: 0.8948
>>> Test Accuracy [200/1000]: 0.9031
>>> Test Accuracy [300/1000]: 0.9074
>>> Test Accuracy [400/1000]: 0.9037
>>> Test Accuracy [500/1000]: 0.9125
>>> Test Accuracy [600/1000]: 0.9135
>>> Test Accuracy [700/1000]: 0.918
>>> Test Accuracy [800/1000]: 0.9152
>>> Test Accuracy [900/1000]: 0.9188
>>> Test Accuracy [999/1000]: 0.9193

```


To run TensorBoard, use the following command (alternatively `python -m tensorflow.tensorboard`) {c} >>tensorboard --logdir=_logs

1.1.1 Cross Entropy on training set by step



title

1.1.2 Accuracy on test set by step (learning curve)



title